



**CIS 412**

**DATABASE MANAGEMENT SYSTEMS**

**Chapter 7**  
**DBMS Functions**

# INTRODUCTION

- Functions of a DBMS
  - Update and retrieve data
  - Provide catalog services
  - Support concurrent update
  - Recover data
  - Provide security services
  - Provide data integrity features
  - Support data independence
  - Support data replication
  - Provide utility services

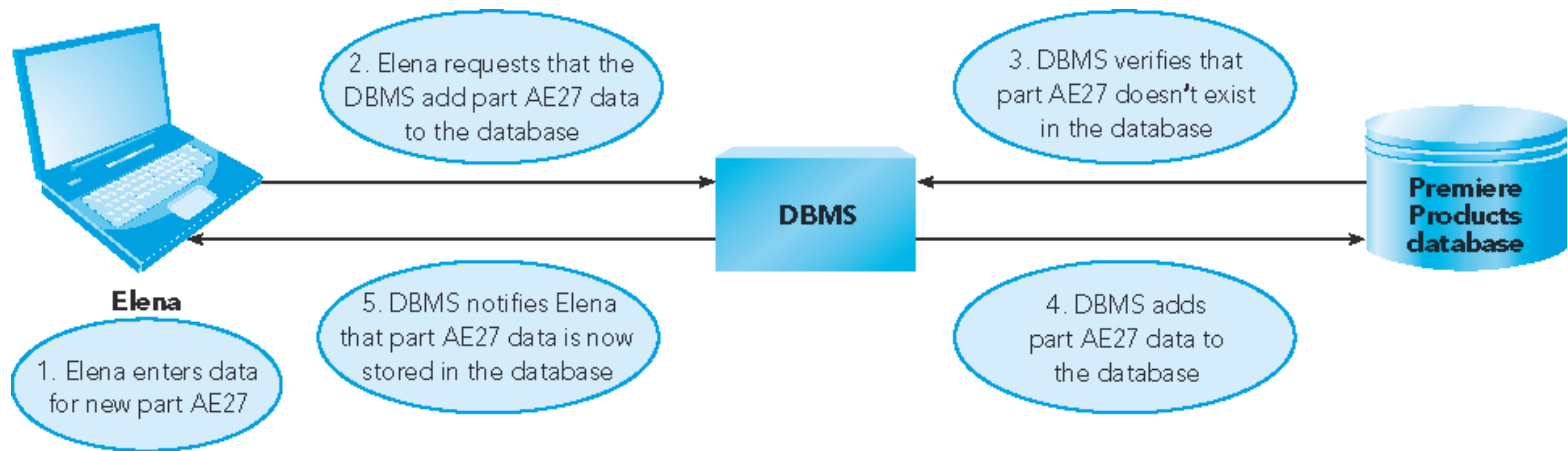


# UPDATE AND RETRIEVE DATA

- Fundamental capability of a DBMS
- Users don't need to know how data is stored or manipulated
- Users add, change, and delete records during updates

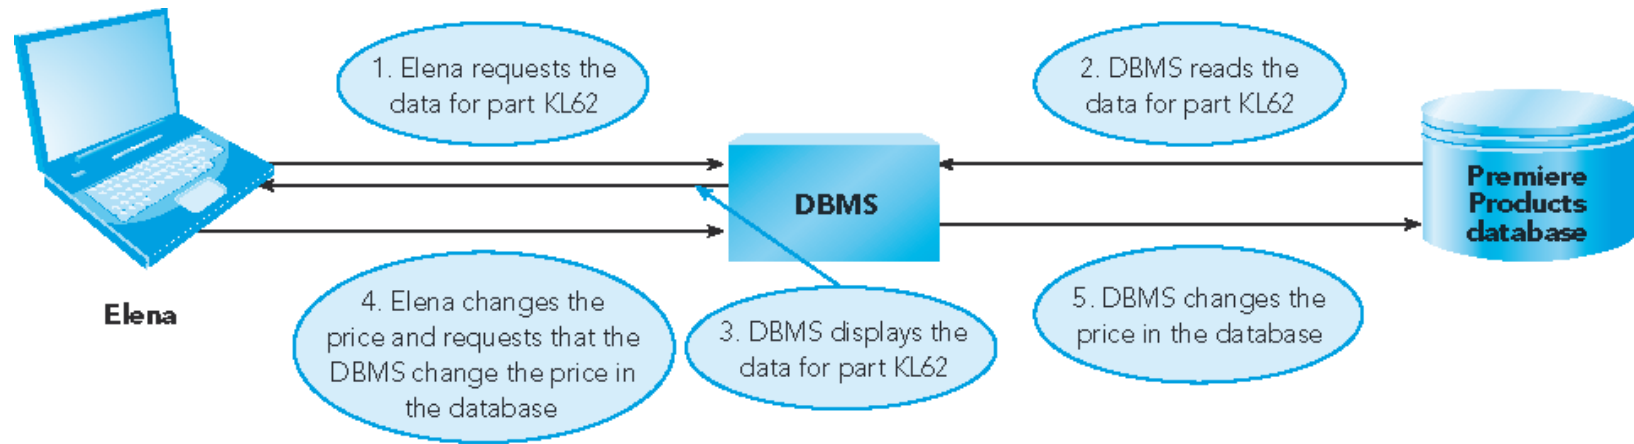


# UPDATE AND RETRIEVE DATA (CONTINUED)




**FIGURE 7-1: Adding a new part to the Premiere Products database**

# UPDATE AND RETRIEVE DATA (CONTINUED)



**FIGURE 7-2: Changing the price of a part in the Premiere Products database**

# UPDATE AND RETRIEVE DATA (CONTINUED)



**FIGURE 7-3: Retrieving a balance amount from the Premiere Products database**

# PROVIDE CATALOG SERVICES

- **Metadata:** data about data
- Stores metadata and makes it accessible to users
- Enterprise DBMSs often have a **data dictionary** (a super catalog)



# SUPPORT CONCURRENT UPDATE

- Ensures accuracy when several users update database at the same time
- Manages complex scenarios for updates
- **Concurrent update:** multiple users make updates to the same database at the same time





# THE CONCURRENT UPDATE PROBLEM

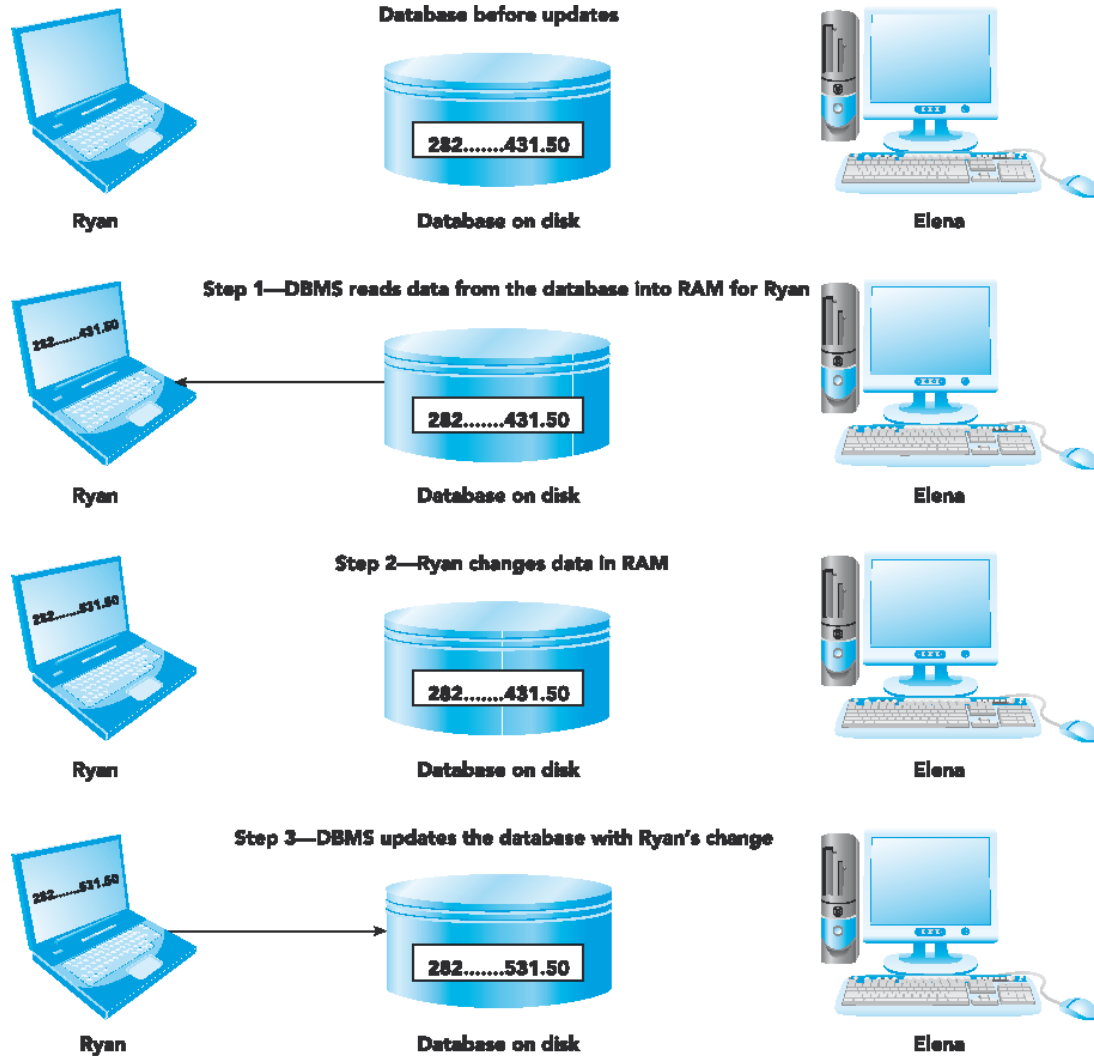
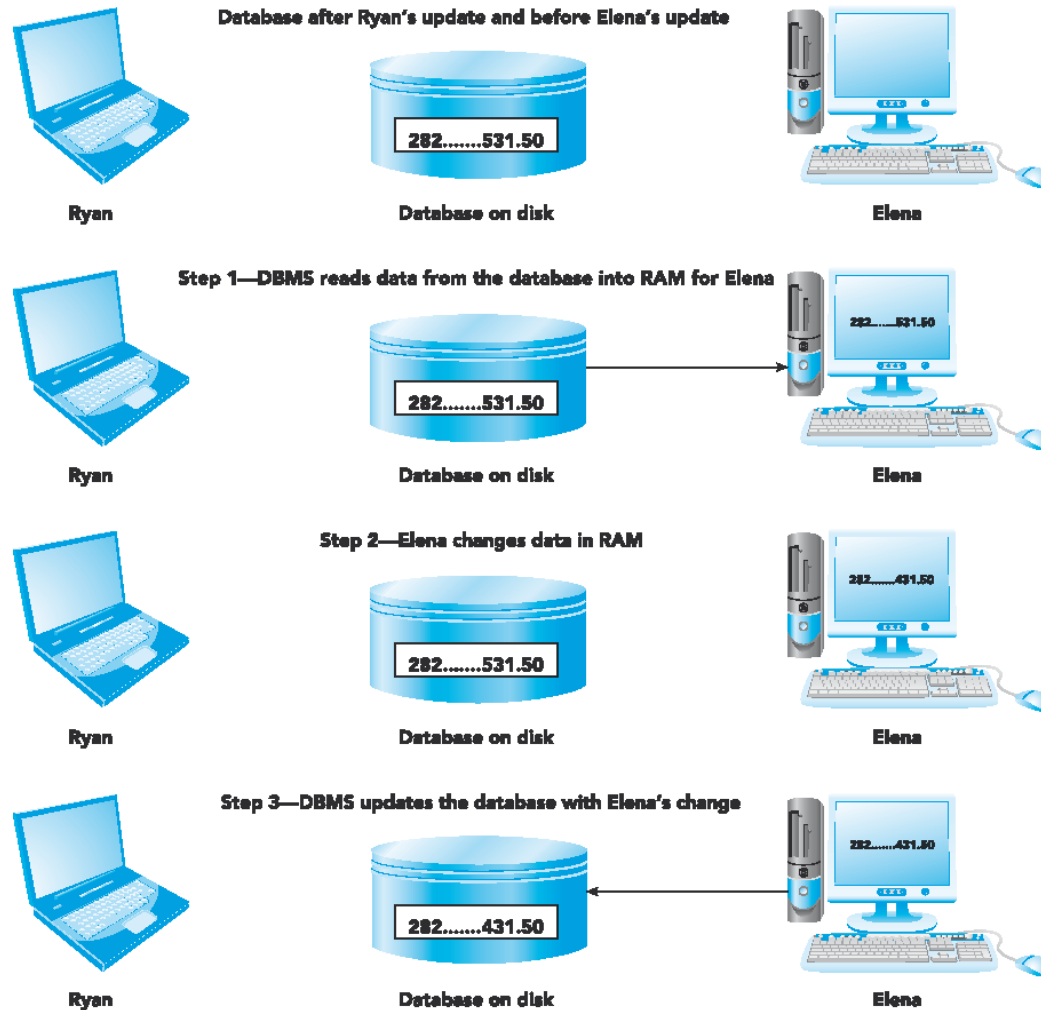


FIGURE 7-4: Ryan updates the database

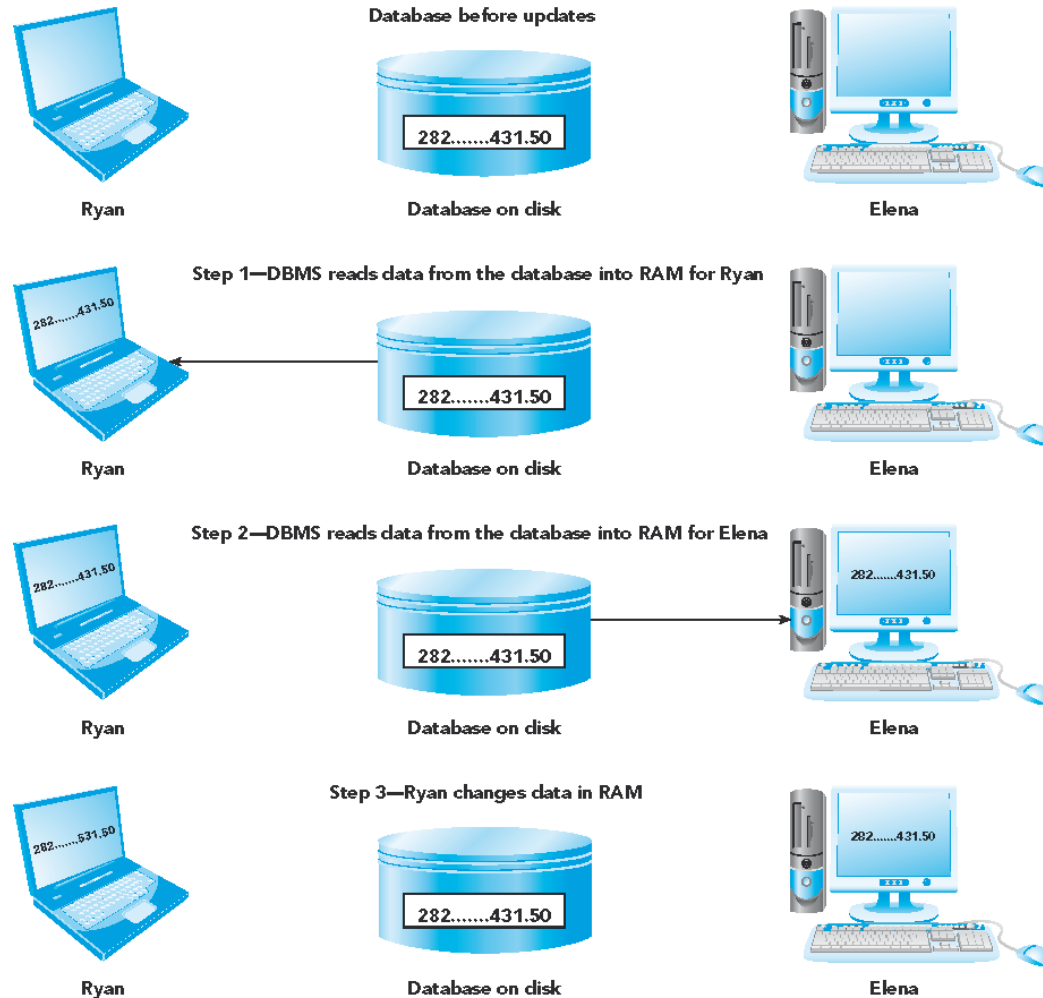


# THE CONCURRENT UPDATE PROBLEM (CONTINUED)



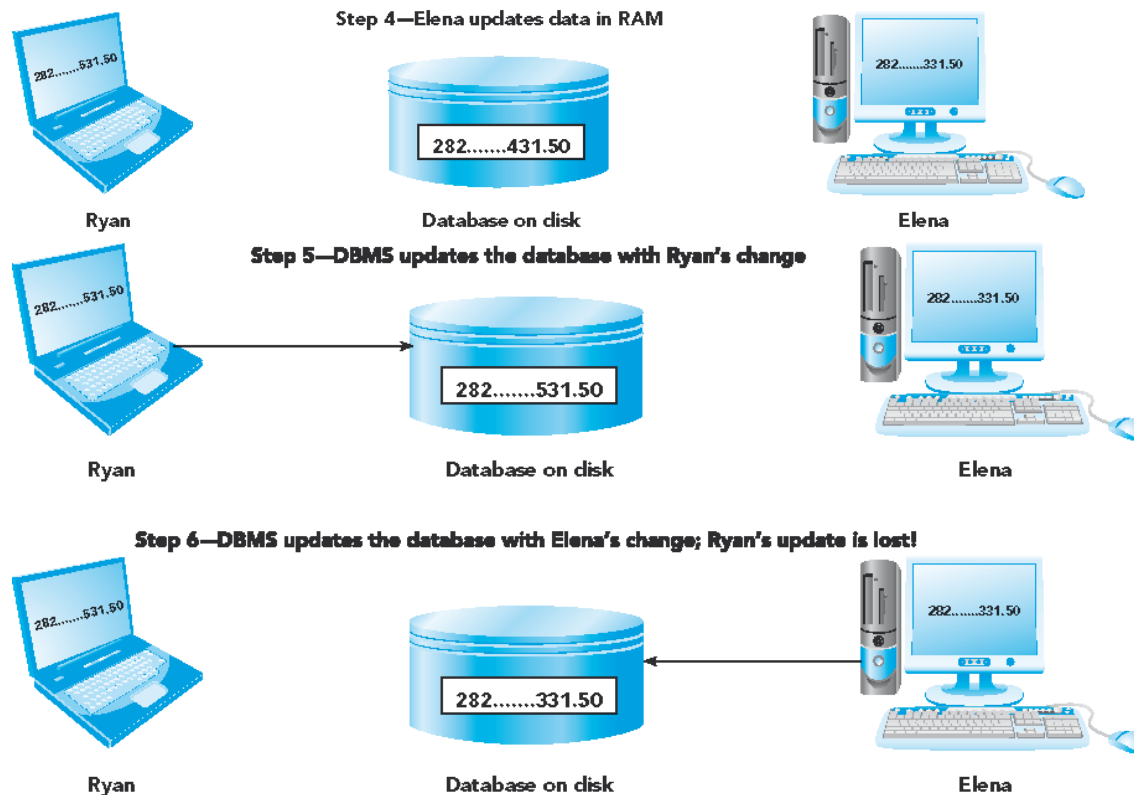
**FIGURE 7-5: Elena updates the database**

# THE CONCURRENT UPDATE PROBLEM (CONTINUED)



**FIGURE 7-6: Ryan's and Elena's updates to the database result in a lost update**

# THE CONCURRENT UPDATE PROBLEM (CONTINUED)



**FIGURE 7-6: Ryan's and Elena's updates to the database result in a lost update**  
(continued)

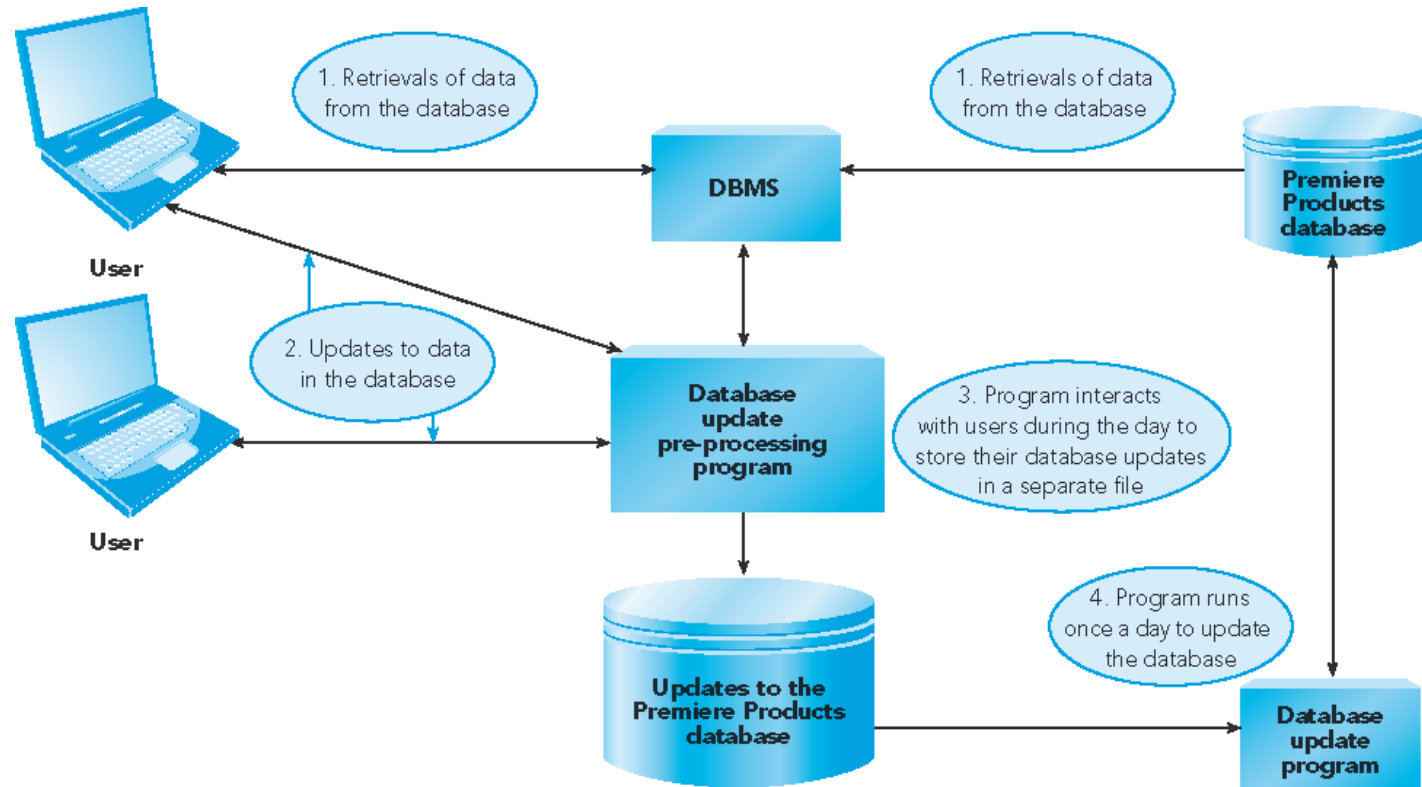
# AVOIDING THE LOST UPDATE PROBLEM

## ○ **Batch processing**

- All updates done through a special program
- Problem: data becomes out of date
- Does not work in situations that require data to be current



# AVOIDING THE LOST UPDATE PROBLEM (CONTINUED)




**FIGURE 7-7: Delaying updates to the Premiere Products database to avoid the lost update problem**

# TWO-PHASE LOCKING

- **Locking:** deny other users access to data while one user's updates are being processed
- **Transaction:** set of steps completed by a DBMS to accomplish a single user task
- **Two-phase locking** solves lost update problem
  - **Growing phase:** DBMS locks more rows and releases none of the locks
  - **Shrinking phase:** DBMS releases all the locks and acquires no new locks



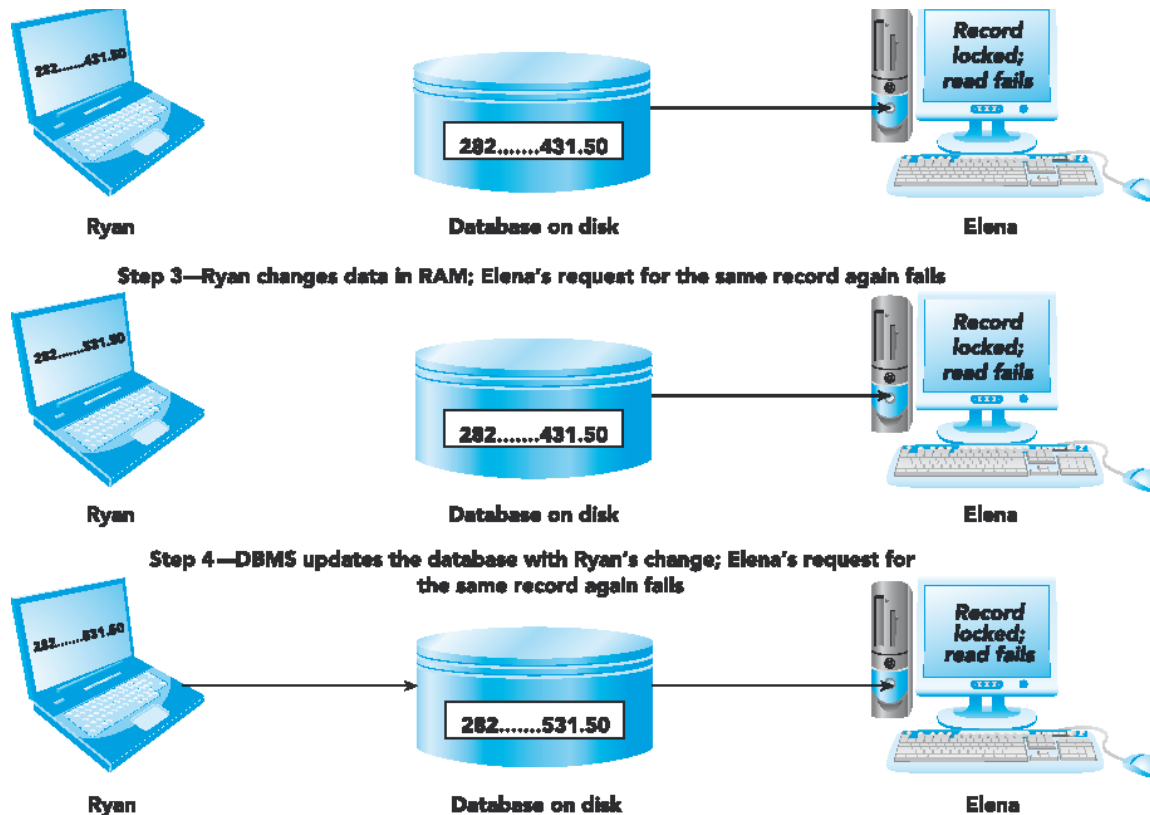
# TWO-PHASE LOCKING (CONTINUED)



**FIGURE 7-8: The DBMS uses a locking scheme to apply Ryan's and Elena's updates to the database**

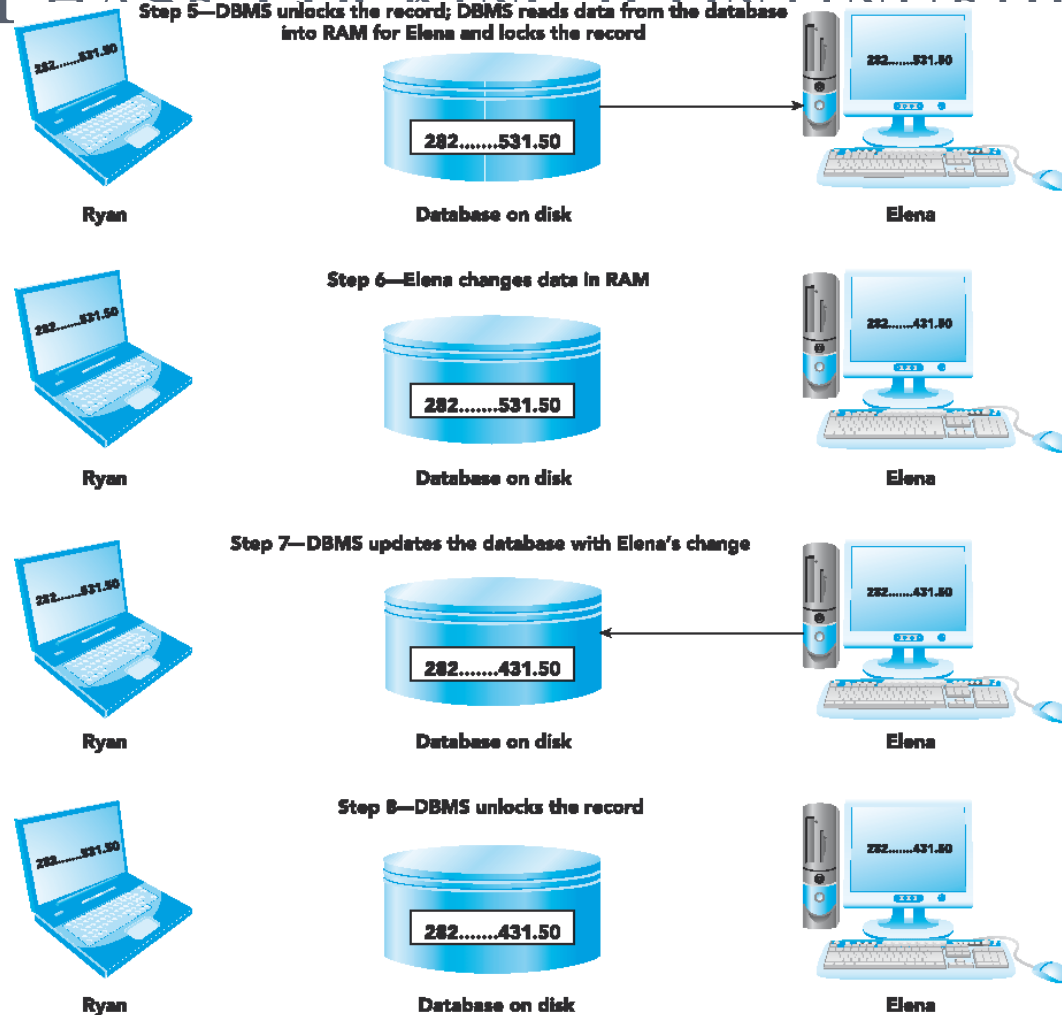


# TWO-PHASE LOCKING (CONTINUED)



**FIGURE 7-8: The DBMS uses a locking scheme to apply Ryan's and Elena's updates to the database (continued)**

# TWO-PHASE LOCKING (CONTINUED)



**FIGURE 7-8: The DBMS uses a locking scheme to apply Ryan's and Elena's updates to the database (continued)**

# DEADLOCK

## ○ **Deadlock or deadly embrace**

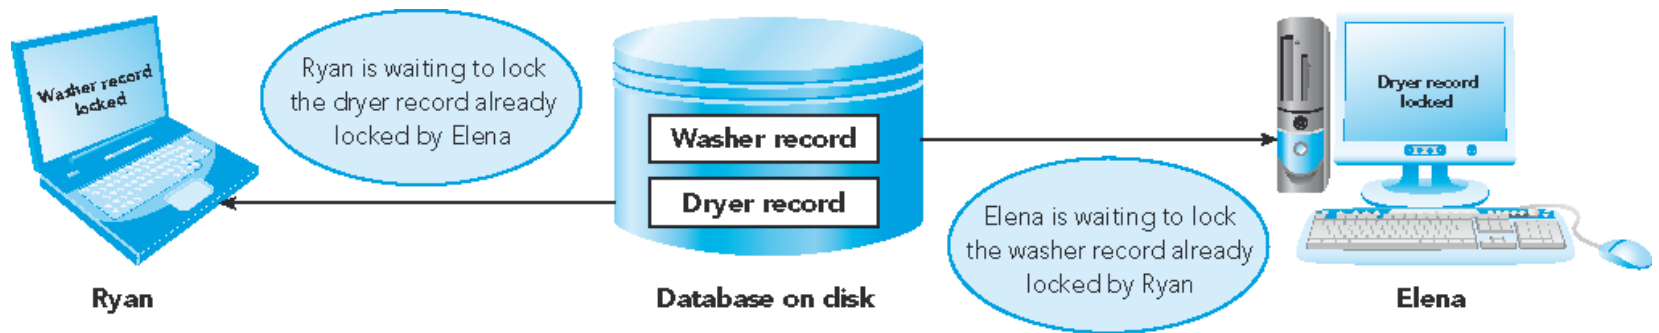
- Two users hold a lock and require a lock on the resource that the other already has
- To minimize occurrence, make sure all programs lock records in the same order whenever possible

## ○ **Managing deadlocks**

- DBMS detects and breaks any deadlock
- DBMS chooses one user to be the **victim**



# DEADLOCK (CONTINUED)



**FIGURE 7-9: Two users experiencing deadlock**



# LOCKING ON PC-BASED DBMSs

- Usually more limited than locking facilities on enterprise DBMSs
- Programs can lock an entire table or an individual row within a table, but only one or the other
- Programs can release any or all of the locks they currently hold
- Programs can inquire whether a given row or table is locked



# TIMESTAMPING

- DBMS assigns each database update a unique time (**timestamp**) when the update started
- Advantages
  - Avoids need to lock rows
  - Eliminates processing time needed to apply and release locks and to detect and resolve deadlocks
- Disadvantages
  - Additional disk and memory space
  - Extra processing time



# RECOVER DATA

- **Recovery:** returning database to a correct state from an incorrect state
- Simplest recovery involves using backups
  - **Backup** or **save:** copy of database



# JOURNALING

- **Journaling**: maintaining a **journal** or **log** of all updates
  - Log is available even if database is destroyed
- Information kept in log for each transaction:
  - Transaction ID
  - Date and time of each update
  - **Before image**
  - **After image**
  - Start of a transaction
  - Successful completion (**commit**) of a transaction





# JOURNALING (CONTINUED)

Transaction ID	Transaction Description
1	1. Change the Price value for part number DW11 to \$389.99
2	1. Add a record to the Orders table: OrderNum of 21700, OrderDate of 10/24/2013, CustomerNum of 282 2. Add a record to the OrderLine table: OrderNum of 21700, PartNum of DW11, NumOrdered of 3, QuotedPrice of \$389.00 3. Add a record to the OrderLine table: OrderNum of 21700, PartNum of KL62, NumOrdered of 2, QuotedPrice of \$346.50 4. Change the OnHand value for part number DW11 to 9 5. Change the OnHand value for part number KL62 to 10 6. Change the Balance value for CustomerNum 282 to \$2,321.50 7. Change the Commission value for RepNum 35 to \$39,346.20
3	1. Add customer 510
4	1. Delete part AT94

**FIGURE 7-10: Four sample transactions**

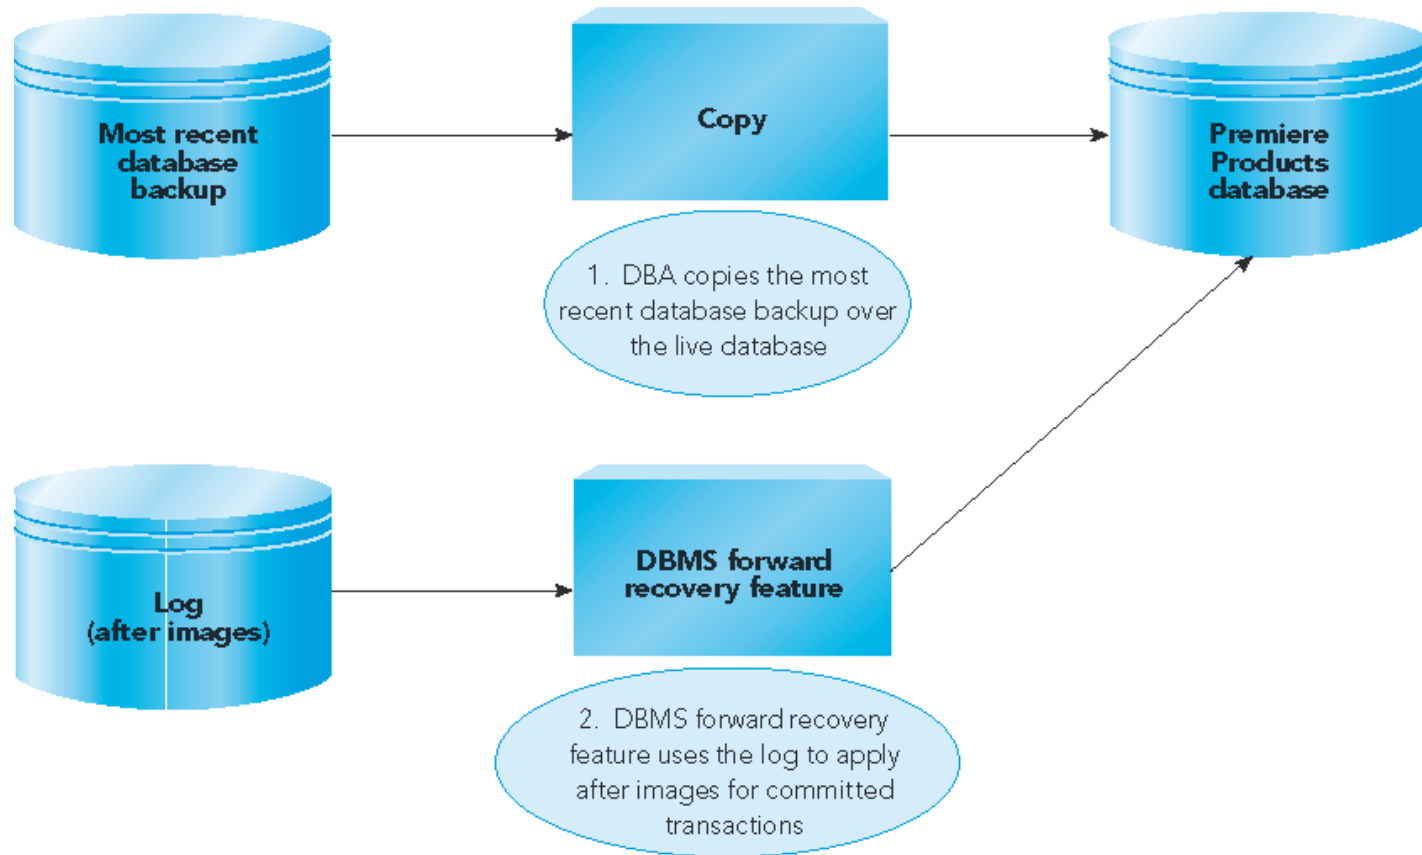


# FORWARD RECOVERY

- DBA executes a DBMS recovery program
- Recovery program applies after images of committed transactions from log to database
- Improving performance of the recovery program
  - Apply the last after image of a record



# FORWARD RECOVERY (CONTINUED)



**FIGURE 7-12: Forward recovery**

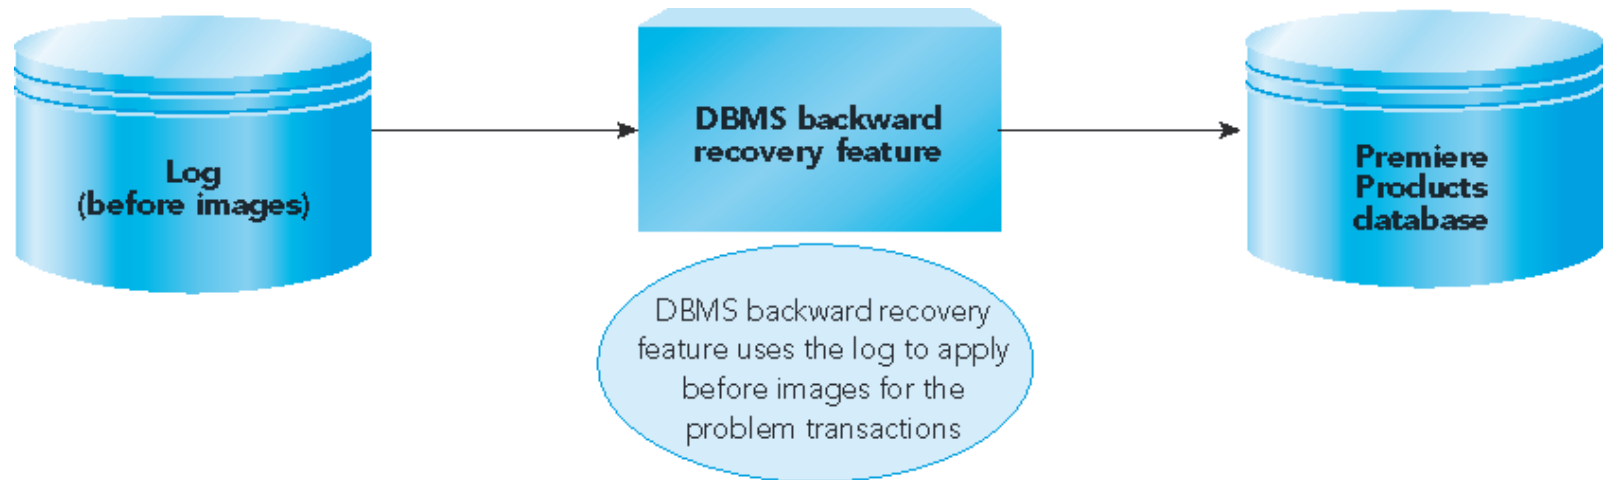


# BACKWARD RECOVERY

- Database not in a valid state
  - Transactions stopped in midstream
  - Incorrect transactions
- **Backward recovery or rollback**
  - Undo problem transactions
  - Apply before images from log to undo their updates



# BACKWARD RECOVERY (CONTINUED)



**FIGURE 7-13: Backward recovery**



# RECOVERY ON PC-BASED DBMSs

- Sophisticated recovery features not available on PC-based DBMSs
- Regularly make backup copies using DBMS
  - Use most recent backup for recovery
- Systems with large number of updates between backups
  - Recovery features not supplied by DBMS need to be included in application programs



# PROVIDE SECURITY SERVICES

- Security: prevention of unauthorized access, either intentional or accidental, to a database
- Most common security features used by DBMSs:
  - Encryption
  - Authentication
  - Authorizations
  - Views



# ENCRYPTION

- **Encryption:** converts data to a format indecipherable to another program and stores it in an encrypted format
- Encryption process is transparent to a legitimate user
- **Decrypting:** reversing the encryption
- In Access, encrypt a database with a password





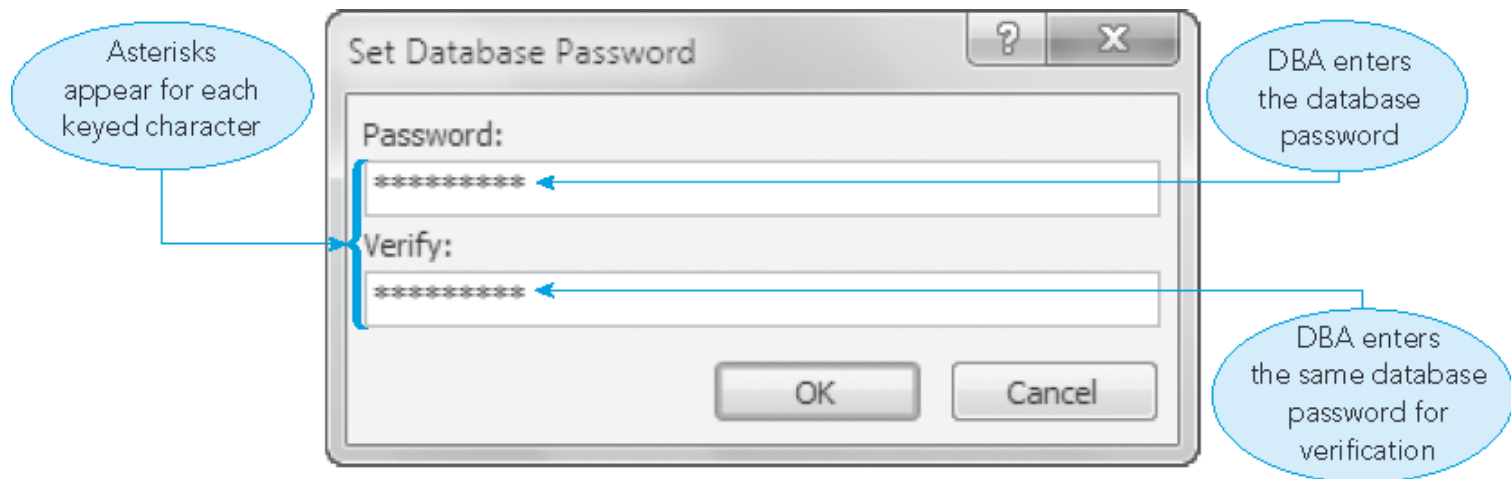
# AUTHENTICATION

- **Authentication:** techniques for identifying the person attempting to access the DBMS
- **Password:** string of characters assigned by DBA to a user that must be entered for access
- **Biometrics:** identify users by physical characteristics such as fingerprints, voiceprints, handwritten signatures, and facial characteristics
- **Smart cards:** small plastic cards with built-in circuits containing processing logic to identify the cardholder



## AUTHENTICATION (CONTINUED)

- **Database password:** string of characters assigned to database that users must enter for accessing the database



**FIGURE 7-14: Assigning a database password to the Premiere Products database**

# AUTHORIZATIONS

- DBA can use **authorization rules** to specify which users have what type of access to which data
- **Permissions**: specify what kind of access the user has to objects in the database
- **Workgroups**: groups of users



# VIEWS

- **View:** snapshot of certain data in the database at a given moment in time
- Can be used for security purposes



# PRIVACY

- **Privacy:** right of individuals to have certain information about them kept confidential
- Laws and regulations dictate some privacy rules
- Companies institute additional privacy rules



# PROVIDE DATA INTEGRITY FEATURES

- Rules followed to ensure data is accurately and consistently updated
- Key integrity
  - Foreign key and primary key constraints
- Data integrity
  - Data type
  - Legal values
  - Format

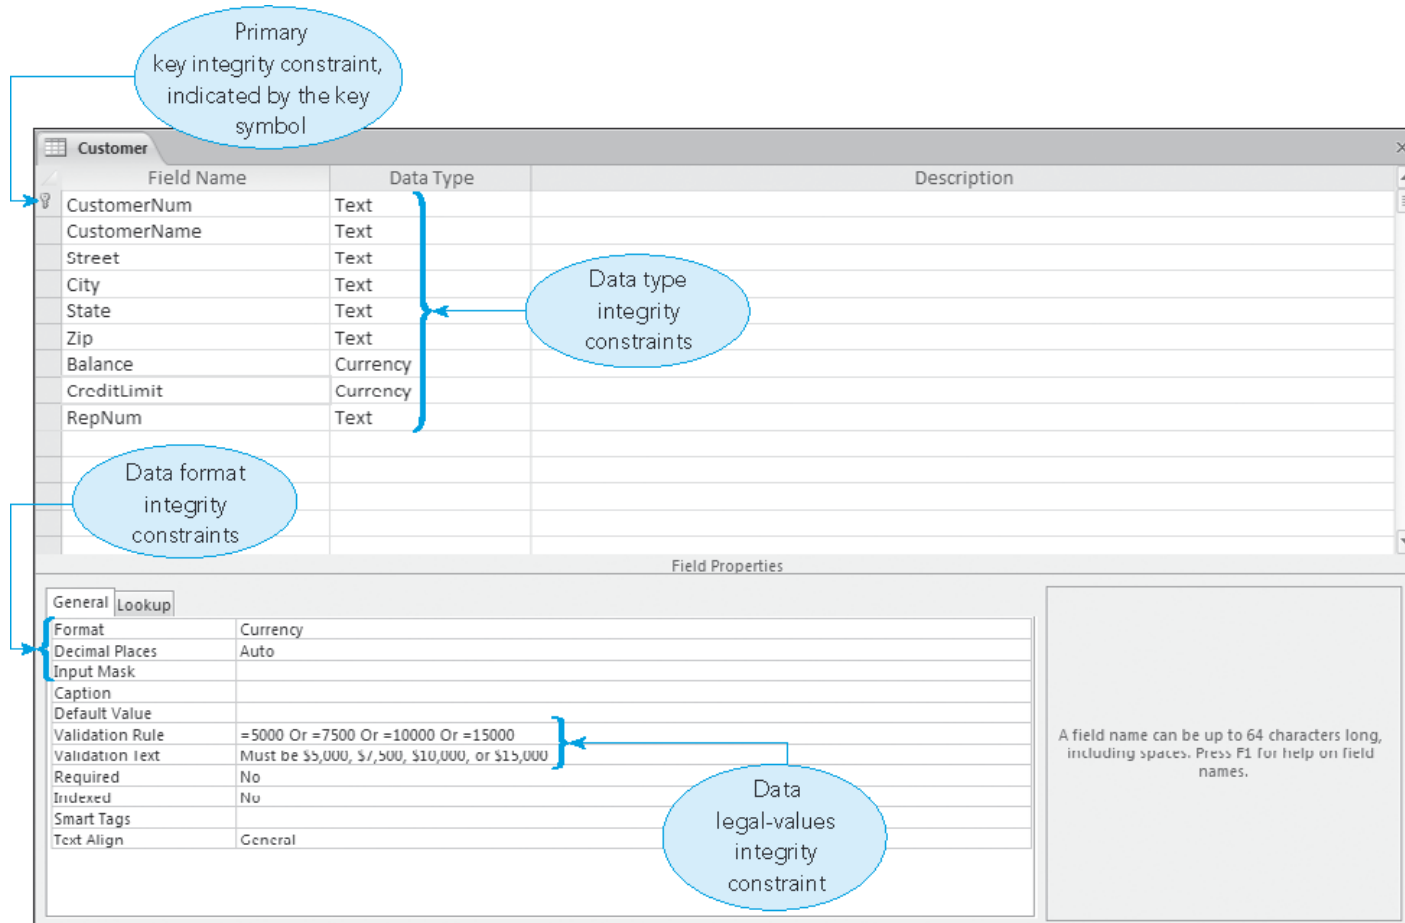


# PROVIDE DATA INTEGRITY FEATURES (CONTINUED)

- Four ways of handling integrity constraints:
  1. Constraint is ignored
  2. Responsibility for constraint enforcement placed on users
  3. Responsibility for constraint enforcement placed on programmers
  4. Responsibility for constraint enforcement placed on DBMS



# PROVIDE DATA INTEGRITY FEATURES (CONTINUED)



**FIGURE 7-16: Example of integrity constraints in Access**



# SUPPORT DATA INDEPENDENCE

- **Data independence:** can change database structure without needing to change programs that access the database
- Types of changes:
  - Adding a field
  - Changing a field property (such as length)
  - Creating an index
  - Adding or changing a relationship



## ADDING A FIELD

- Don't need to change any program except those programs using the new field
- SQL SELECT \* FROM command will present an extra field
  - Solution: list the required fields in an SQL SELECT command instead of using \*



# CHANGING THE LENGTH OF A FIELD

- Generally, don't need to change programs
- Need to change the program if:
  - Certain portion of screen or report is set aside for the field and the space cannot fit the new length



# CREATING AN INDEX

- To create an index, enter a simple SQL command or select a few options
- Most DBMSs use the new index automatically
- For some DBMSs, need to make minor changes in already existing programs



# ADDING OR CHANGING A RELATIONSHIP

- Trickiest of all
- May need to restructure database




# SUPPORT DATA REPLICATION

- **Replicated:** duplicated
- Manage multiple copies of same data in multiple locations
- Maintained for performance or other reasons
- Ease of access and portability
- **Replicas:** copies
- **Synchronization:** DBMS exchanges all updated data between master database and a replica



# SUPPORT DATA REPLICATION (CONTINUED)



**FIGURE 7-18: DBMS synchronizes two databases in a replica set**



# PROVIDE UTILITY SERVICES

- **Utility services** assist in general database maintenance
- Change database structure
- Add new indexes and delete indexes
- Use services available from operating system
- Export and import data
- Support for easy-to-use edit and query capabilities, screen generators, report generators, etc.





## PROVIDE UTILITY SERVICES (CONTINUED)

- Support for procedural and nonprocedural languages
  - **Procedural language:** must tell computer precisely how a given task is to be accomplished
  - **Nonprocedural language:** describe task you want computer to accomplish
- Easy-to-use menu-driven or switchboard-driven interface



# SUMMARY

- DBMS allows users to update and retrieve data in a database without needing to know how data is structured on disk or manipulated
- DBMS must store metadata (data about the data) and make this data accessible to users
- DBMS must support concurrent update
- Locking denies access by other users to data while DBMS processes one user's updates
- During *deadlock* and *deadly embrace*, two or more users are waiting for the other user to release a lock before they can proceed



## SUMMARY (CONTINUED)

- In timestamping, DBMS processes updates to a database in timestamp order
- DBMS must provide methods to recover a database in the event the database is damaged
- DBMSs provide facilities for periodically making a backup copy of the database
- Enterprise DBMSs maintain a log or journal of all database updates since the last backup; log is used in recovery process



## SUMMARY (CONTINUED)

- DBMSs provide security features (encryption, authentication, authorizations, and views) to prevent unauthorized access to a database
- DBMS must follow rules or integrity constraints (key integrity constraints and data integrity constraints) so that it updates data accurately and consistently
- DBMS must support data independence
- DBMS must have facility to handle data replication
- DBMS must provide utility services that assist in general maintenance of a database

